CONFIDENTIAL

Clever Investor Platform

Product & Org Proposal

Prepared for Luke Babich & Ben Mizes · March 2026

Varun Raghavan

STRUCTURE

What this proposal covers

01
Situation
Where Clever Offers is today — what mystery shopping revealed

02
The Gap
5 UX gaps, structural conflict of interest, competitive benchmarks

03
The Hypothesis
What the Investor Platform could be — and why it's the right bet

04
Product Proposal
Scrappiest path to validation, consumer experience, data schema + AI

05
Org & Resources
Team constraints, what to protect, how to staff and sequence

06
Open Questions
What I'd want to explore with you in this session

Key assumptions & how I prepared

What I did

- Mystery shopped cleveroffers.com end-to-end
- Spoke with 2 Clever-connected agents (Luke Smith, Jacob Mkhitarian) — 17 and 12 min calls
- Mystery shopped InvestorLift and HomeLight Simple Sale
- Reviewed Luke's async answers on team structure, growth signals, and investor relationships

What I'm assuming

- No formal market research — the thesis is grounded in Clever's first-hand experience with sellers, investors, and agents
- Hiring timelines of 4–12 weeks are realistic, not optimistic
- AMP is P1 and fully protected — no pull from that team
- Flat hiring cost by role type (PM ~$160K, Eng ~$180K, fully loaded)
- ATLAS enterprise fire drills are unpredictable — treated as unavailable

What I'd still want to know

- Conversion rate from submit → investor offer accepted (current baseline)
- Who owns the early investor-matching work and what's been built
- Whether any investor relationships have yielded repeat transactions

SITUATION — CURRENT STATE

What the seller experience actually looks like today

Homepage

Zip → See Cash Offers
Property + address
Low friction, clean UX

Contact gate

Name/email/phone required before anything shown

"Offers on the way"

No timeline. No dashboard. No status.

Agent calls

Not investors. No context. No account.

In-person visit

Required before any investor sees property

⚠️ **Conversion chokepoint**

Contact required before value shown

⚠️ **Dead zone**

No persistence. No seller account.

⚠️ **Channel conflict**

Agent incentive is to convert to listing

⚠️ **Visit bottleneck**

Adds 1–2 weeks before investor sees property

First-hand finding: Both mystery-shopped agents spent the first 2 minutes on investor options, then 15+ minutes pitching traditional listings. The platform promises cash offers — sellers get agent calls, then silence. There is no Clever account to land on after submission.

COMPETITIVE CONTEXT

The market has already solved pieces of this

WHAT CLEVER OWNS

Clever Offers

- ✓ Trusted consumer brand (4.9★, 3,400+ reviews)
- ✓ 18,000+ agent network, validated seller funnel
- ✓ 30% of revenue — growing 250% YoY
- ✗ No buy-box matching, no investor platform
- ✗ Experience ends at submission

Consumer trust — missing the backend

BEST CONSUMER COMP

HomeLight Simple Sale

- ✓ Named investor network surfaced to seller
- ✓ Explicit investor vs. agent path choice
- ✓ Structured intake call feeds pricing model
- ✓ No in-person visit before investor sees property
- ✗ Smaller agent network than Clever

Right consumer model — weaker network

BEST INVESTOR-SIDE COMP

InvestorLift

- ✓ AI buy-box matching — investor sets criteria
- ✓ 5M+ investors, $30B+ offers made
- ✓ AI deal scoring + in-app negotiation
- ✓ Live marketplace visible before sign-up
- ✗ B2B only — no consumer-safe experience

Best matching intelligence — no B2C layer

THE HYPOTHESIS

Two bets — and the second only works because of the first.

HYPOTHESIS 1

The consumer relationship ends at submission.

Today, Clever hands off and disappears. A persistent account layer — starting with how the seller schedules their concierge call — is the missing foundation every downstream feature requires. The same account extends beyond Offers: it benefits every consumer persona Clever serves.

HYPOTHESIS 2

Clever can build the investor marketplace neither HomeLight nor InvestorLift can.

Clever already has the consumer trust and agent network both competitors lack. Building InvestorLift-style matching intelligence behind Clever's existing experience creates a marketplace positioned to win — but only if a consumer account layer exists to qualify and inform the sellers being routed.

These aren't separate bets. H1 is what makes H2 work — an investor marketplace without a consumer account layer is a better blast list, not a platform.

C Consumer trust

Clever already owns the seller relationship

H HomeLight model

Parallel paths, named network, structured intake

IL InvestorLift intelligence

Buy-box matching, AI scoring, in-app negotiation

= Clever Investor Platform · InvestorLift-style backend · Clever's consumer front-end · account layer that compounds across every product line

PRODUCT PROPOSAL — REC 1

Rec 1: Replace the dead zone with a scheduling moment

Current — what sellers receive today

- ✗ "Cash offers are on the way!" — no timeline
- ✗ CX assigns agent (seller unaware)
- ✗ Unclear voicemail + cold text introduction
- ✗ Two agent intros, no context on who or why
- ✗ No Clever account to land on

→

Proposed: seller-scheduled

- ✓ **Channel selection**
  Phone or SMS
- ✓ **Quick-time chips**
  5 / 10 / 15 min / later today
- ✓ **Date + time picker**
  Specific scheduling
- ✓ **Concierge pre-briefed**
  Preferences sent ahead of call
- ✓ **Account created**
  First persistent identity layer

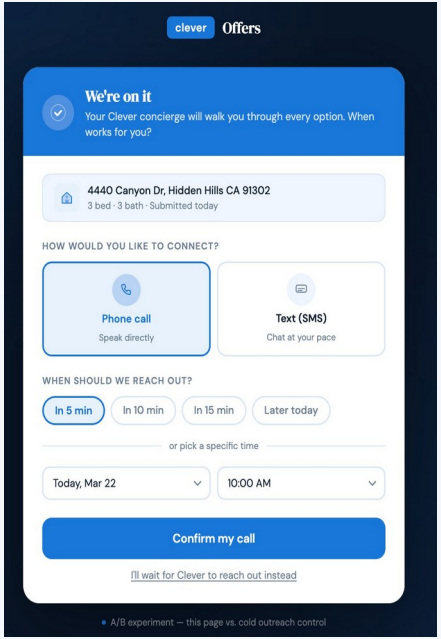

A/B Experiment: Control (cold outreach) vs. Test A (scheduled · phone only) vs. Test B (scheduled · channel choice) — primary metric: submit → concierge connection rate

PRODUCT PROPOSAL — REC 2

Rec 2: Give sellers a place to land and a reason to engage

Personalized home card

Address, submitted date, property details — the seller sees their home, not a generic confirmation

5-step progress tracker

Submitted → Call scheduled → Explore options → Get offers → Decide. Seller always knows where they are

Local market snapshot

Median price, avg DOM, list-to-sale ratio for their zip. Educates before the concierge call

Interactive path comparison

Cash offer vs. open market — side-by-side with local estimates. Seller selects preference

Preference module → agent brief

Priority, timeline, showing comfort, reason for selling. Selections generate a structured brief the concierge inherits

Foundation, not a feature — creates the first persistent Clever consumer account. Every future feature — for sellers, buyers, and every persona we serve — lives here.

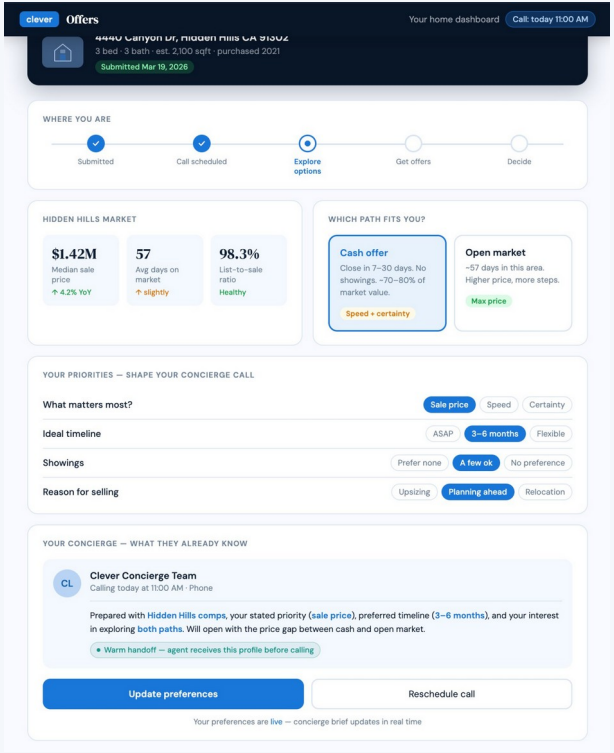

STRATEGIC CONTEXT

This isn't a Clever Offers feature.

It's the account layer that makes Clever a platform.

Clever's business model ends at the referral across every product line. The post-submit account layer — seller identity, preferences, market context — has direct compounding value beyond Offers.

CORE BUSINESS

Agent Matching

Seller preferences (timeline, priority, path) make every agent match warmer. Same data schema, same infrastructure — zero additional build required.

Now: Cold lead → agent call with no context

CLOSING SERVICES

Clever Title

A seller with a dashboard and progress tracker is the right place to surface title and escrow mid-journey — at Get Offers, not as a cold pitch at close.

Now: Title introduced at close — unfamiliar, cold, last minute

LENDER CHANNEL

Clever Pro

Seller account data — mortgage balance, timeline, reason for selling — is pre-qualification signal. Raw leads become enriched pipeline for lender partners.

Now: Raw contact passed to lender with no intent data

LONG-TERM LOYALTY

Gravy / Lifecycle

The Gravy acquisition was about owning the pre-purchase journey. A seller dashboard is the same thesis on the sell side — holding the relationship across the lifecycle.

Now: Clever exits after referral — no post-transaction relationship

PRODUCT PROPOSAL — INTELLIGENCE LAYER

Seller preferences become structured data — LLM turns data into a brief

seller_profile

seller_id · address · property_type

beds · baths · est_sqft · purchase_year

improvements[ ] · est_mortgage_balance

channel_preference · concierge_scheduled_at

seller_preferences

path_preference (investor | open_market | undecided)

primary_priority · ideal_timeline

showing_comfort · reason_for_selling

path_card_clicks[ ] — behavioral signal

time_on_page · sections_viewed[ ]

market_context

zip · median_sale_price · avg_days_on_market

list_to_sale_ratio · yoy_price_change

investor_activity_score

comparable_cash_offer_range[ ]

open_market_estimate_range[ ]

↓ LLM

Agent brief (LLM output):

"Varun is a 3bd/3ba homeowner in Hidden Hills who purchased in 2021 and has made meaningful improvements. He's planning ahead for aging parents — no urgency. He leaned toward cash offers but listed sale price as his top priority, which is a tension worth surfacing early. Local market: median $1.42M, 57 DOM. Cash offer range: $980K–$1.15M vs. open market $1.35M–$1.48M. Lead with the price gap and let him decide."

The scrappiest path to validating the investor matching hypothesis

Phase 1

Month 1–2 · Loveable

Structured intake + manual routing

→ Agent collects property package via structured Loveable form (photos, details, timeline)

→ Clever platform routes to curated investor shortlist — no agent personal lists

→ Investor responses tracked in Airtable/Sheets

→ Validation metric: do structured packages generate more investor responses than ad-hoc agent outreach?

Phase 2

Month 3–4 · +ANVIL

Buy-box matching + algorithmic routing

→ ANVIL builds investor buy-box scoring — adjacent to existing match models

→ Property intake auto-routes to investors whose buy-box criteria match

→ Seller dashboard surfaces: 'X investors reviewing your home'

→ Validation metric: does buy-box matching increase offer volume per submission?

Phase 3

Month 5+ · Full build

Competing offers dashboard

→ Seller sees competing offers forming in real time — side by side with open market estimate

→ In-platform negotiation layer (vs. agent-mediated verbal relay)

→ Agent shifts from bottleneck → advisor: in the loop on offers, not generating them

→ Validation metric: submit → offer acceptance rate vs. control

Key constraint: Phase 1 and 2 require no tech allocation from AMP, ATLAS, or Agent HQ. Phase 3 requires ANVIL partial and a dedicated new pod.

RESOURCE ALLOCATION

What to protect, what to borrow, what to build

AMP	PROTECT	P1 — fully staffed. Zero pull. This is the engine. Do not touch.
ATLAS	PROTECT	Too stretched + enterprise fire drill risk. Unavailable for net-new work.
Agent HQ	PROTECT	Understaffed due to attrition. Fix the backfill first. Don't add scope.
ANVIL	PHASE 2 ONLY	Named as intended Phase 2 partner for buy-box scoring — adjacent to existing match models. Not committed until capacity confirmed in live session.
Growth	1 SPRINT	Funnel instrumentation for the A/B experiment only. One shared sprint with Ben's team.
NEW POD	BUILD	Mandate: Offers investor matching + post-submit account layer (cross-product). 1 PM, Month 1. Engineers gated on Phase 1 signal: 15%+ investor response lift OR 3 investor matches in 30 days.

Phase 1 requires zero pull from any existing pod — Loveable + dedicated PM only. Engineering gated on signal: 15%+ investor response lift OR 3 investor matches in 30 days. ANVIL capacity and new pod reporting structure are open questions for the live session.

PRINCIPLES & MEASUREMENT

How we make decisions — and how we know it's working

01 The consumer relationship is the asset

A consumer who trusts Clever is a consumer we can stay with throughout their transaction — beyond the match. That continuity creates more value for them, for Clever, and new monetization opportunities.

02 Build platforms to invest in the long-term consumer relationship

Every capability we ship should deepen our ability to serve the consumer over time — a data schema, a signal, an experience layer that compounds across every product line.

03 Signal before scale

Demonstrating demand leads to full investment. Speed to learning is a feature, and we're always pointed towards learning something new — whether that signal confirms the bet or redirects it.

04 Agents are allies, not obstacles

Clever's agent network is a competitive moat. We design investor matching to run parallel to agents — not in place of them.

OKRs — illustrative

O1: Make Clever the most trusted way for any consumer to navigate their options

- KR1: Seller-to-concierge connection rate (submit → live conversation)
- KR2: Consumer account activation — persistent Clever accounts beyond initial submission
- KR3: Consumer NPS on Offers path at offer receipt and post-decision

O2: Build the investor matching foundation that makes Clever's intelligence a durable competitive advantage

- KR1: Investor-to-match conversion on structured vs. unstructured leads
- KR2: Referral revenue from Offers path, growing Q-o-Q from current signal baseline

Targets set collaboratively with the team once baselines are established in week one.

RECOMMENDATION

Four bets, in order.

1 Rec 1: Ship the scheduling screen now

It's Loveable-executable, requires no tech allocation, and directly addresses the highest-friction point in the current flow. It also creates the first seller account — the foundation for everything else.

2 Rec 2: Build the seller dashboard in parallel

The preference module generates the data schema. The LLM turns it into agent briefs. The A/B experiment runs itself. This is the concierge upgrade that makes warm handoffs systematic — not dependent on individual agent quality.

3 Start Phase 1 investor intake with Loveable

Replace the agent personal-list blast with a structured property package and a platform-owned investor routing layer. No new tech team required for Month 1. Signal within 6–8 weeks.

4 Gate Phase 2 on Phase 1 signal

Only pull ANVIL and fund the new pod when Phase 1 shows offer volume improvement. Don't build the matching engine on a hypothesis — build it on evidence.

LIVE SESSION AGENDA

What I'd want to dig into with you

The early investor-matching work

10 min

Who owns it, what's been built, what's worked and what hasn't. I want to build on what exists, not past it.

New pod sequencing

10 min

Does this start under Ben's team for Phase 1 and graduate at Phase 2, or does it need dedicated ownership from Month 1?

Pressure-test the Phase 1 scope

10 min

Does the structured intake + Loveable routing proposal fit the Offers team's actual capacity? What am I missing about current dependencies?

What's been tried and didn't work

10 min

What has Clever already tested on Offers that didn't land? That signal is as valuable as what did.

ANVIL capacity and the Phase 2 gate

10 min

Is ANVIL available for Phase 2 buy-box scoring, or does that need to be an early engineering hire? This changes how I'd sequence the pod build.

Reactions and open curveballs

10 min

Where am I wrong? What assumptions are off? What's the thing you're not sure I understand yet about how Clever actually works?